

Elasticsearch-1.6.0

监控&插件&源码调试

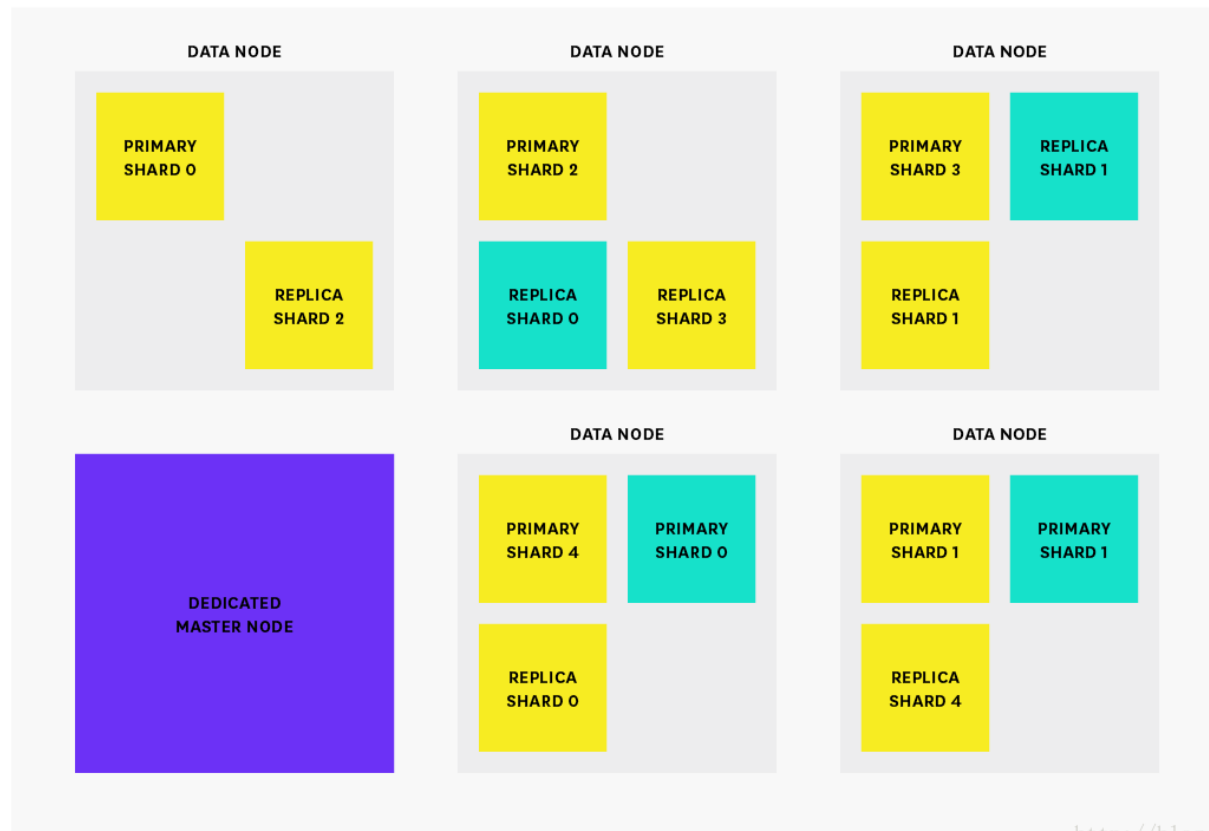
王志

2017.12

1. Elasticsearch脚本监控
2. Elasticsearch 插件介绍
3. Elasticsearch 源码调试

Elasticsearch 监控

Cluster



Legend



- elasticsearch cluster (集群整体状态)
- Elasticsearch node (集群节点状态)

Elasticsearch监控

- Cluster的监控 (localhost:9200)
 1. [Cluster Health](#) - 查看集群健康状态接口 (/_cluster/health)
 2. [Cluster State](#) - 查看集群状况接口 (/_cluster/state)
 3. [Cluster Stats](#) - 查看集群统计信息接口 (/_cluster/stats)
 4. [Pending cluster tasks](#) - 查看集群任务 (_cluster/pending_tasks)
- Nodes层面的监控 (localhost:9200)
 1. [Nodes Stats](#) - 节点状态 (/_nodes/stats)
 2. [Nodes Info](#) - 节点信息 (/_nodes)
 3. [Nodes hot_threads](#) - 节点的热线程(/_nodes/hot_threads)

Elasticsearch监控

- **Elasticsearch 性能监控指标**

1. Search and indexing performance - 搜索性能指标
2. Memory and garbage collection - 索引性能指标
3. Host-level system and network metrics - 内存使用和GC指标
4. Cluster health and node availability - 集群健康和节点可用性
5. Resource saturation and errors - elasticsearch主机的网络 and 系统
6. Resource saturation and errors - 资源saturation and errors

参考链接：<http://www.jianshu.com/p/6574e2288745>

Elasticsearch 监控

• 搜索性能指标的要点

1. Query load : 监控当前请求数量
2. Query latency: query阶段时延 (总耗时/总查询次数)
3. Fetch latency: 提取阶段平均时延

Metric description	Name	[Metric type]
Total number of queries	indices.search.query_total	Work: Throughput
Total time spent on queries	indices.search.query_time_in_millis	Work: Performance
Number of queries currently in progress	indices.search.query_current	Work: Throughput
Total number of fetches	indices.search.fetch_total	Work: Throughput
Total time spent on fetches	indices.search.fetch_time_in_millis	Work: Performance
Number of fetches currently in progress	indices.search.fetch_current	Work: Throughput



Elasticsearch 监控查询

• 索引性能指标

1. Indexing latency :监控index的平均耗时
2. Refresh latency : 监控refresh的平均耗时
3. Flush latency : 监控flush的平均耗时

Metric description	Name	[Metric type]
Total number of documents indexed	indices.indexing.index_total	Work: Throughput
Total time spent indexing documents	indices.indexing.index_time_in_millis	Work: Performance
Number of documents currently being indexed	indices.indexing.index_current	Work: Throughput
Total number of index refreshes	indices.refresh.total	Work: Throughput
Total time spent refreshing indices	indices.refresh.total_time_in_millis	Work: Performance
Total number of index flushes to disk	indices.flush.total	Work: Throughput
Total time spent on flushing indices to disk	indices.flush.total_time_in_millis	Work: Performance



Elasticsearch 监控查询

• 内存使用和GC指标

1. Garbage collection duration and frequency
2. JVM heap in use、JVM heap used、JVM heap committed、Memory usage

Metric description	Name	[Metric type]
Total count of young-generation garbage collections	jvm.gc.collectors.young.collection_count	Other
Total time spent on young-generation garbage collections	jvm.gc.collectors.young.collection_time_in_millis	Other
Total count of old-generation garbage collections	jvm.gc.collectors.old.collection_count	Other
Total time spent on old-generation garbage collections	jvm.gc.collectors.old.collection_time_in_millis	Other
Percent of JVM heap currently in use	jvm.mem.heap_used_percent	Resource: Utilization
Amount of JVM heap committed	jvm.mem.heap_committed_in_bytes	Resource: Utilization

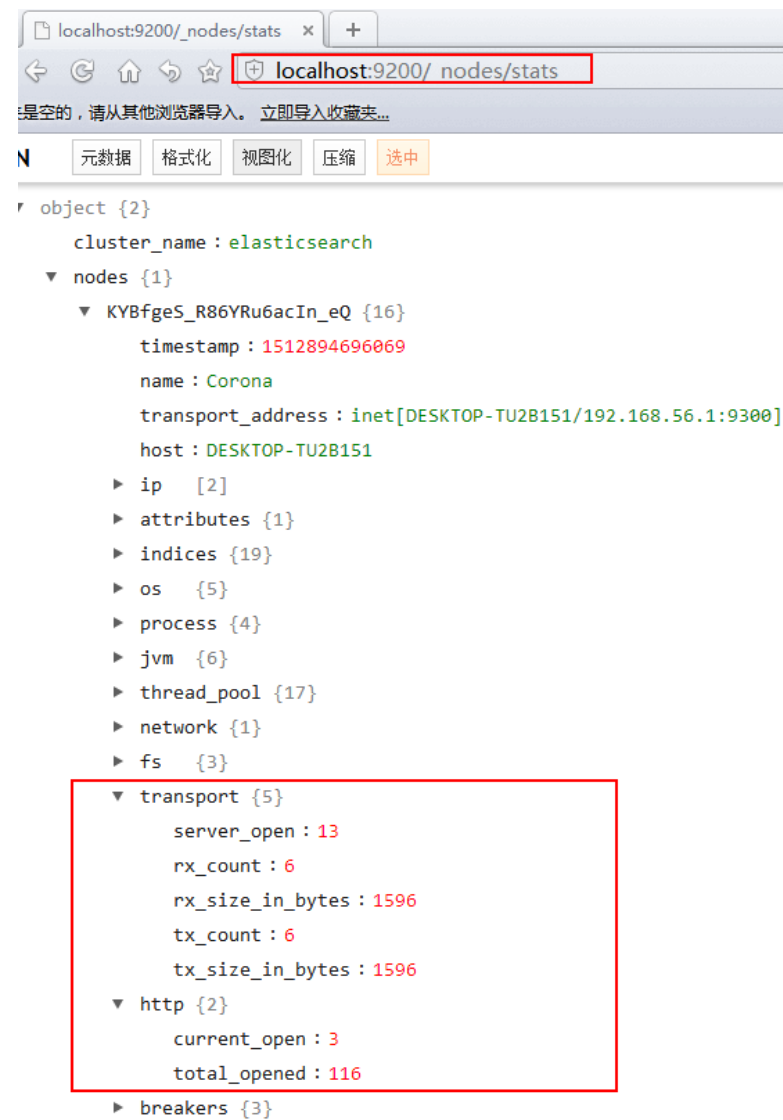


Elasticsearch 监控查询

• elasticsearch主机的网络和系统

1. Disk space
2. I/O utilization
3. CPU utilization
4. Network bytes sent/received
5. Open file descriptors
6. HTTP connections

Name	[Metric type]	
Available disk space	Resource: Utilization	
I/O utilization	Resource: Utilization	
CPU usage	Resource: Utilization	
Network bytes sent/received	Resource: Utilization	
Open file descriptors	Resource: Utilization	
Metric description	Name	[Metric type]
Number of HTTP connections currently open	http.current_open	Resource: Utilization
Total number of HTTP connections opened over time	http.total_opened	Resource: Utilization

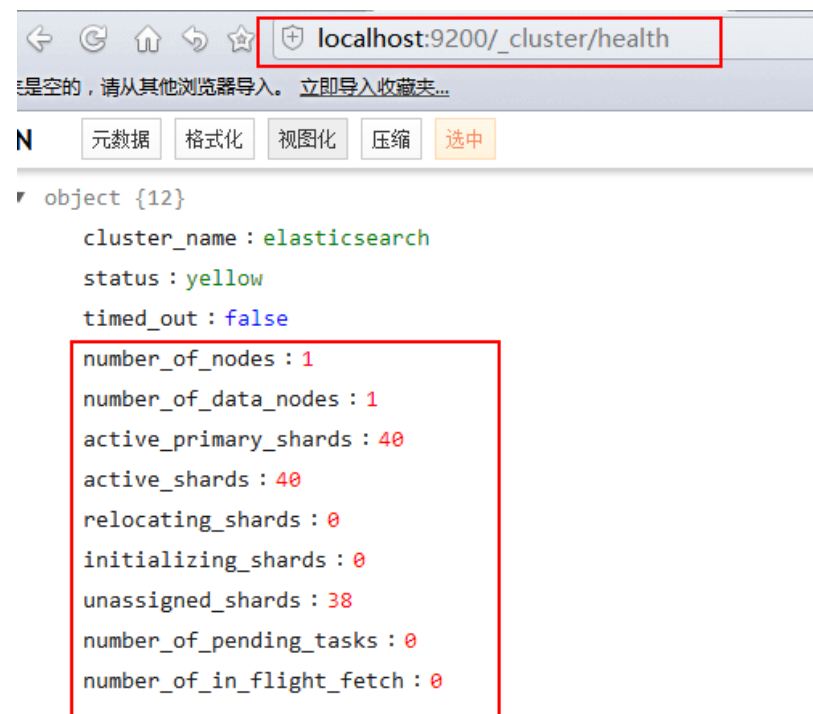


Elasticsearch 监控查询

集群健康和节点可用性

1. Cluster status: 黄色代表副本分片未分配或丢失，红色代表主分片丢失。
2. Initializing and unassigned shards：分片仍处于初始化或未分配状态

Metric description	Name	[Metric type]
Cluster status (green, yellow, red)	cluster.health.status	Other
Number of nodes	cluster.health.number_of_nodes	Resource: Availability
Number of initializing shards	cluster.health.initializing_shards	Resource: Availability
Number of unassigned shards	cluster.health.unassigned_shards	Resource: Availability



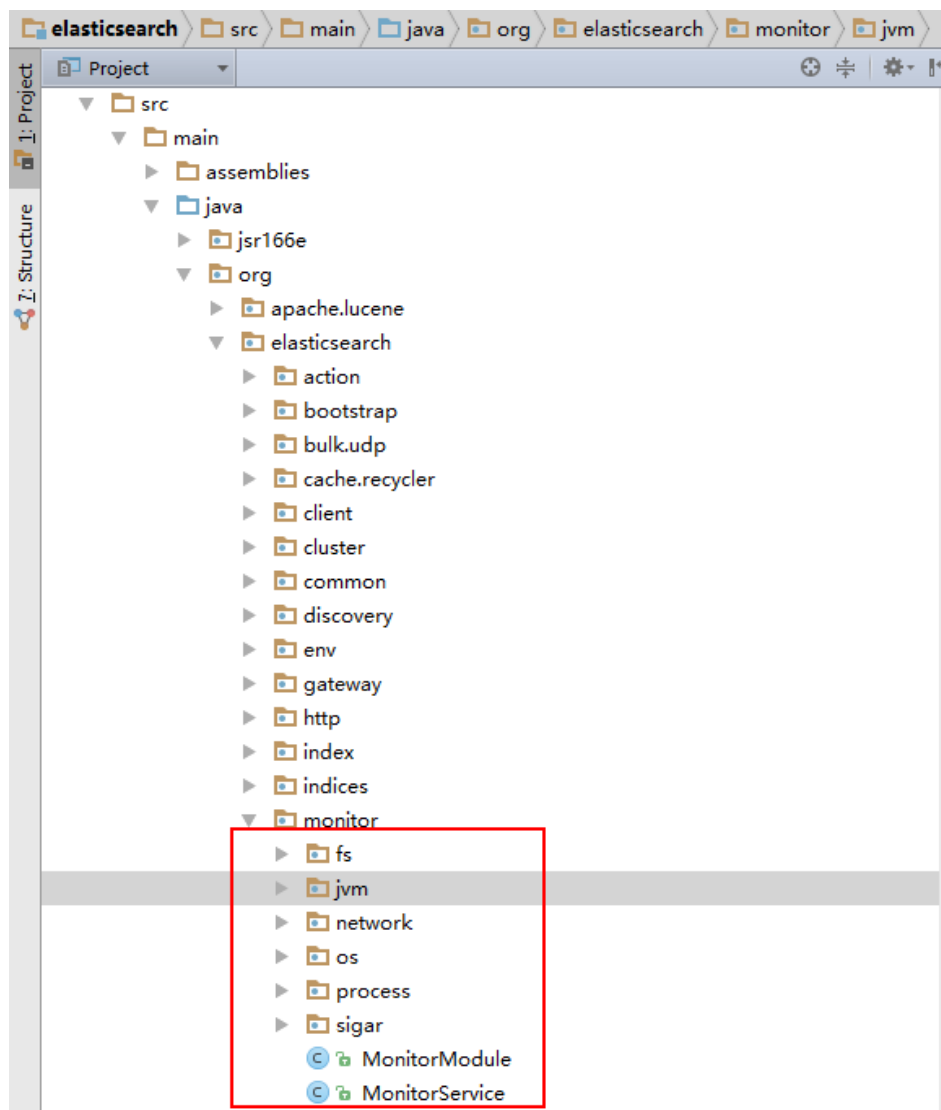
Elasticsearch 监控查询

- 监控源码探究

1. JMX
2. Sigar

- 监控数据采集

1. 文档见附件
2. Demo : [github](#)

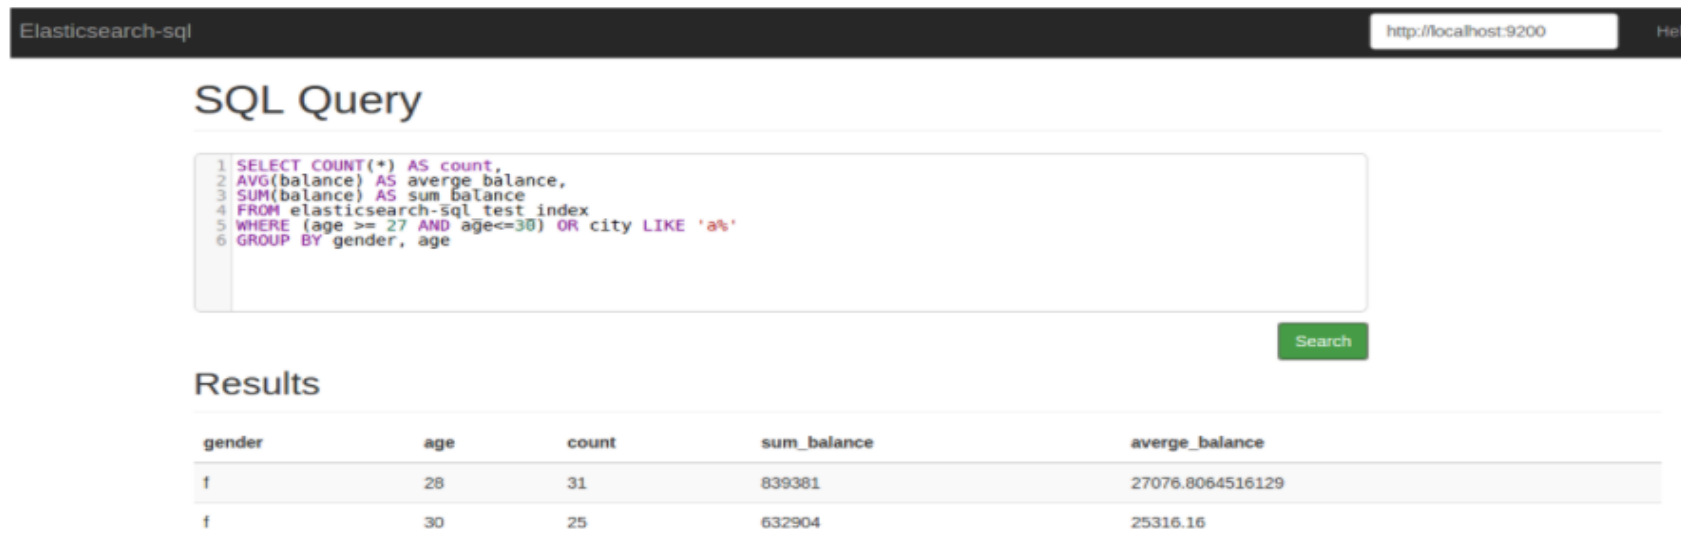


1. Elasticsearch脚本监控
2. Elasticsearch 插件介绍
3. Elasticsearch 源码调试

Elasticsearch 插件介绍

- 插件介绍

1. Head
2. Bigdeck
3. Sql
4. more



Elasticsearch 插件介绍

• 插件安装

1. 自动安装：

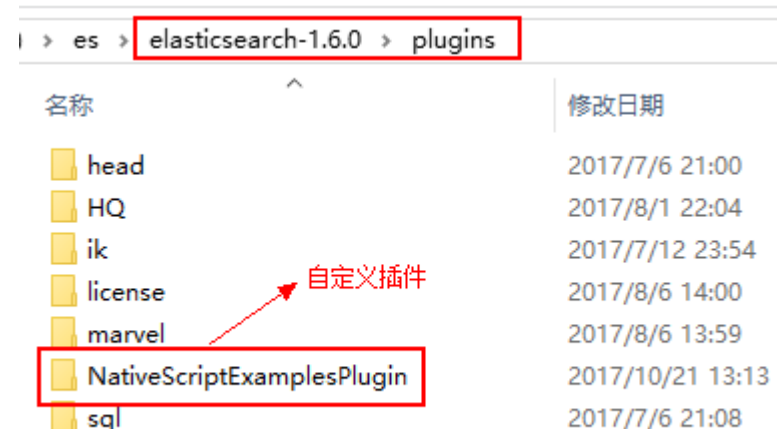
1. 在Elasticsearch目录下
2. `$/bin/plugin -install mobz/elasticsearch-head`

2. 手动安装

1. 在elasticsearch-1.6.0\plugins新建目录
2. 拷贝自定义插件的jar包到目录下
3. Es重启自动加载插件

3. 插件查看

1. `http://localhost:9200/_cat/plugins`



Elasticsearch 插件介绍

- 插件开发
 1. 打分(script)插件：用于自定义打分
 2. Rest插件：用于浏览器访问
 3. Cluster插件：用于操作集群所有节点
- Demo
 - https://github.com/lebron374/elasticsearch_cluster_operation
- 参考资料（doc文档为插件加载过程）
 - <http://www.jianshu.com/p/b32b838d5220>
 - <http://www.jianshu.com/p/17d4477c0954>



Elasticsearch 插件介绍

- 插件开发
 - 定义Plugin插件类
 - 定义打分script类并通过plugin类注册
 - 新增es-es-plugin.properties
 - 打包编译生成jar包
- 参考资料
 - 参考文档：<http://www.jianshu.com/p/7a9af40a29f7>
 - Demo：https://github.com/lebron374/elasticsearch_cluster_operation
 - Demo：https://github.com/lebron374/elasticsearch_rest_plugin_demo

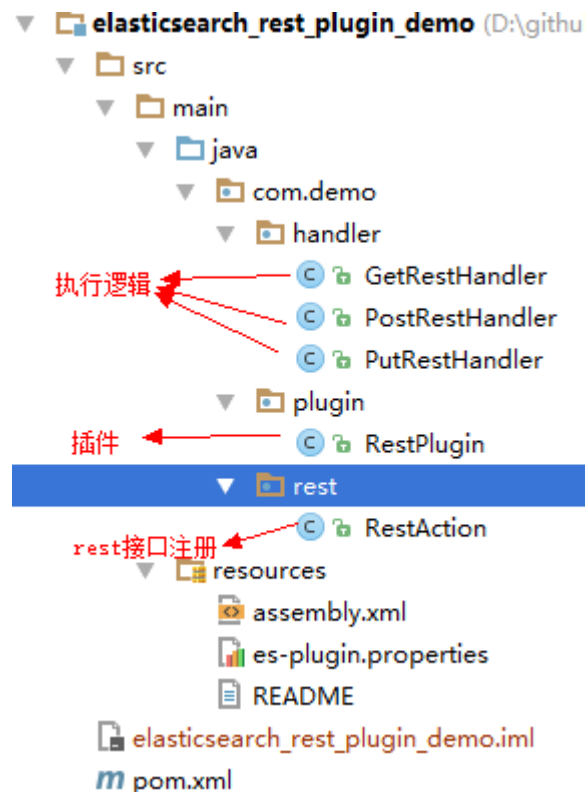
Elasticsearch - 打分插件

The screenshot illustrates the development of an Elasticsearch plugin named 'elasticsearch_cluster_operation'. The project structure is shown in the top-left pane, with annotations identifying key components: 'HeroPlugin' as the plugin class, 'HeroScript' as the scoring script, and 'es-plugin.properties' as the plugin definition file.

The top-right pane displays the Java code for 'HeroScript' and 'HeroPlugin'. 'HeroScript' implements 'NativeScriptFactory' and defines a scoring script named 'hero_script'. 'HeroPlugin' extends 'AbstractPlugin' and registers the 'HeroScript' plugin.

The bottom pane shows the 'es-plugin.properties' file, which defines the plugin's name and the class that implements the 'ScriptEngine' interface. The line 'plugin=com.demo.plugin.HeroPlugin' is highlighted, with an annotation '指定插件类' (Specify plugin class) pointing to it.

Elasticsearch - rest插件



```
public class RestAction extends BaseRestHandler {  
  
    @Inject  
    public RestAction(Settings settings, Client client, RestController controller) {  
        super(settings, controller, client);  
  
        controller.registerHandler(RestRequest.Method.GET, "/rest", new GetRestHandler(client));  
        controller.registerHandler(RestRequest.Method.POST, "/rest", new PostRestHandler(client));  
        controller.registerHandler(RestRequest.Method.PUT, "/rest", new PutRestHandler(client));  
    }  
  
    @Override  
    protected void handleRequest(RestRequest request, RestChannel channel, Client client) throws Exception {  
    }  
}
```

注册接口

```
public class RestPlugin extends AbstractPlugin {  
  
    @Override  
    public String name() { return "rest demo"; }  
  
    @Override  
    public String description() { return "rest demo"; }  
  
    /**  
     * @param module  
     */  
    public void onModule(RestModule module) { module.addAction(RestAction.class); }  
}
```

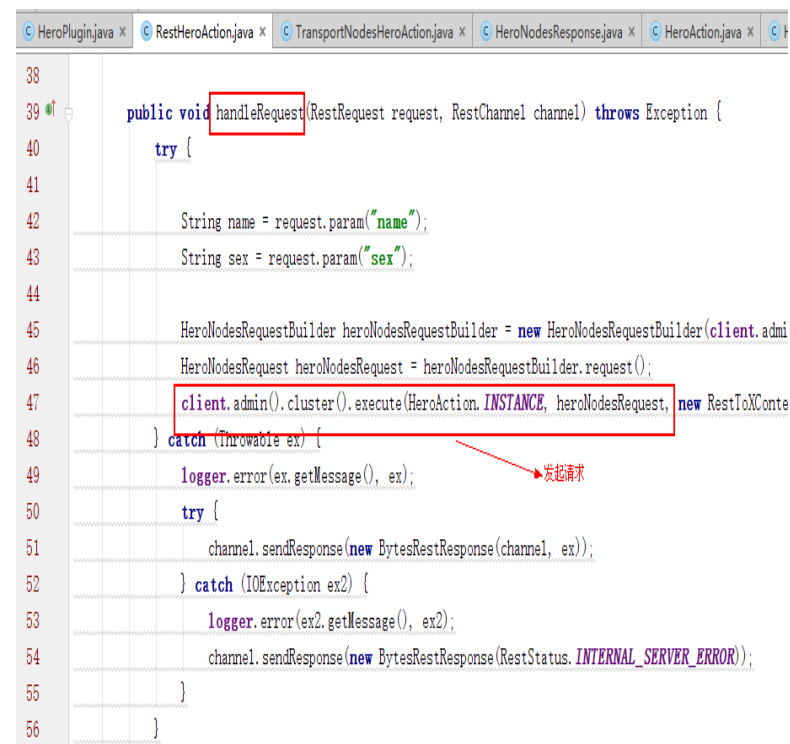
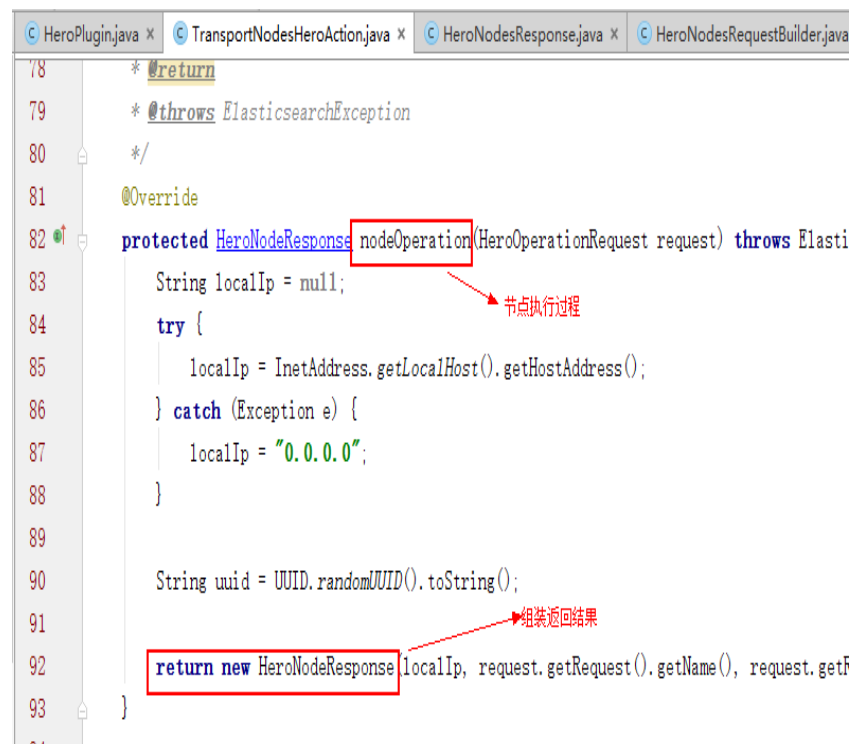
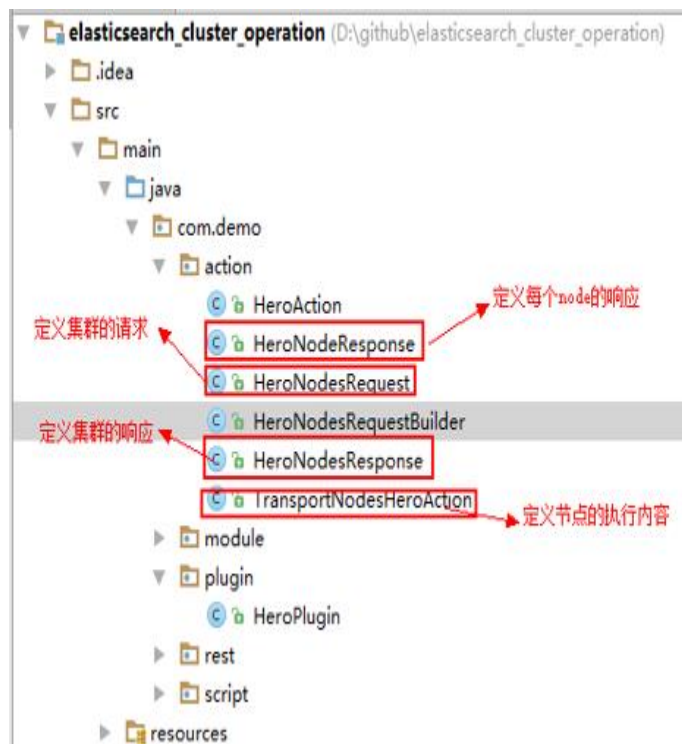
注册插件

```
public class GetRestHandler implements RestHandler {  
  
    private final ESLogger esLogger = ESLoggerFactory.getLogger(GetRestHandler.class.getName());  
  
    private Client client;  
    private String prefix;  
  
    public GetRestHandler(Client client) { this.client = client; }  
  
    @Override  
    public void handleRequest(RestRequest request, RestChannel channel) throws Exception {  
        /* 获取参数的方法 */  
        prefix = request.param("prefix");  
        /* 执行具体的动作 */  
        NodesInfoResponse response = client.admin().cluster().prepareNodesInfo().all().execute().actionGet();  
        List<String> nodes = new ArrayList<>();  
        for (NodeInfo nodeInfo : response.getNodes()) {  
            String nodeName = nodeInfo.getNode().getName();  
            nodes.add(nodeName);  
        }  
  
        /* 回传数据 */  
        try {  
            sendResponse(request, channel, nodes);  
        } catch (IOException ioe) {  
            esLogger.error("Error sending response", ioe);  
        }  
    }  
  
    private void sendResponse(RestRequest request, RestChannel channel, List nodes) throws IOException {  
    }  
}
```

处理请求

返回响应

Elasticsearch – cluster插件



1. Elasticsearch脚本监控
2. Elasticsearch 插件介绍
3. Elasticsearch 源码调试

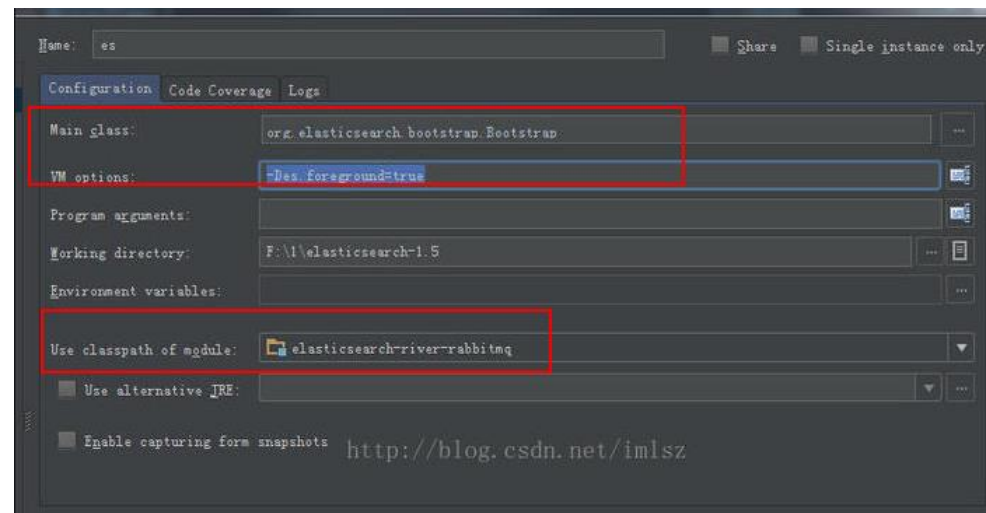
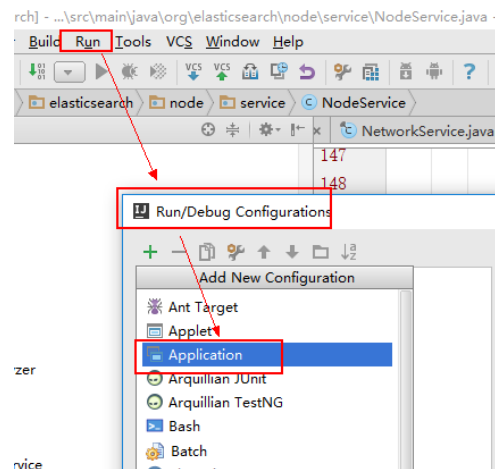
Elasticsearch 源码调试

- 源码调试步骤

1. 导入elasticsearch源码
2. 导入自定义插件的源码
3. Idea上按照右图进行配置

- 参考资料

<http://www.jianshu.com/p/83d49e3e341b>



Elasticsearch 源码参考

- 参考资料

1. 源码阅读：<http://www.jianshu.com/p/e84f868335b3>
2. 插件介绍：<https://www.cnblogs.com/huangfox/p/3541300.html>
3. 自定义插件：<http://www.jianshu.com/p/7a9af40a29f7>
4. Plugin demo：https://github.com/lebron374/elasticsearch_cluster_operation
https://github.com/lebron374/elasticsearch_rest_plugin_demo
5. Es性能监控：<http://www.jianshu.com/p/6574e2288745>